

Lecture 6: Parallel computing, cloud computing and working on Amazon Web Services

Greg Caporaso

gregcaporaso@gmail.com

Some last thoughts on regular expressions

Robust searches

- Sometimes your queries will fail
 - Won't produce output (good)
 - Will produce incorrect output (bad)
- Fail loudly! Produce a (useful) error message on failure.


Designing robust searches

- Make assumptions explicit
 - If you're assuming that your records start with '>', search for '^>' to avoid matching '>' characters that show up in other places
- Match full lines by including ^ and \$ in your search query
- Check the number of matches that were made: is it reasonable?

Testing of software

- Start thinking about what positive and negative controls for these terms might look like. Software testing is something we'll be discussing regularly through-out the semester.

Why is parallel computing important in bioinformatics?



Cluster computing

- Many computers connected to one another to serve as a larger compute resource.
- Compute-intensive jobs can be split over many systems and run in parallel.
- Similar to desktop compute hardware, but different casing, no (or only few) displays/keyboards directly connected.
- Owned and maintained “in-house”.

Why is parallel computing important in bioinformatics?

| Platform | Sanger | 454 (Titanium) | Illumina Genome Analyzer II | Illumina HiSeq2000 | Illumina MiSeq |
|---|-----------|----------------|-----------------------------|--------------------------|-------------------------------|
| Read Length (bases) | ~1000 | ~400 | 150 (single end) | 100 (single end) | 150 (single end); 250 soon |
| Number of reads | 96 or 384 | ~1,000,000 | ~100,000,000 | ~1,600,000,000 | ~10,000,000 |
| Maximum number of samples per run | n/a | 1000 | 12,000 (barcode-limited) | 24,000 (barcode-limited) | 2500 (barcode-limited) |
| Sequences per \$1 (sequencing costs only) | 0.44 | 100 | 5000 | 200,000 | 12,500 |

The “benchtop” sequencer



MiSeq Instrument

The MiSeq system is a fully integrated sequencing ecosystem, in a compact and economical instrument. For results in hours, not days, MiSeq uses TruSeq, Illumina's reversible terminator-based sequencing by synthesis chemistry to deliver the fastest time to answer. Perform the widest breadth of sequencing applications, including highly multiplexed PCR amplicon sequencing, small genome resequencing and de novo sequencing, small RNA sequencing, library quality control, and 16S metagenomics, with automated, on-instrument data analysis workflows to take your research further.

[Do you have questions?](#)

What's New

[Sequencing Portfolio Brochure](#)

09/29/2011

[MiSeq Brochure](#)

09/29/2011

[MiSeq System Product Information Sheet](#)

09/29/2011

Overview

Applications

Featured Researchers

Performance and Specifications >

Technology

Scientific Data

Workflow

Kits

Approximate Run Duration and Output

Cluster Generation and Sequencing

| Read Length | Total Time for Amplification and Sequencing* | Output** |
|-------------|--|------------|
| 1 × 36 bp | ~4 hours | 175-245 Mb |
| 2 × 25 bp | ~5 hours | 250-350 Mb |
| 2 × 100 bp | ~19 hours | 1.0-1.4 Gb |
| 2 × 150 bp | ~27 hours | 1.5-2.0 Gb |

*Includes paired-end read, if applicable.

** Install specifications for MiSeq with an Illumina PhiX library and cluster densities between 720-880 k/mm² that pass filtering. Performance may vary based on sample quality, cluster density, and other experimental factors.


OTU picking: example of a compute intensive process

What taxa are represented in each sample?

```
>PC.634_1 FLP3FBN01ELBSX
CTGGGCCGTGTCTCAGTCCCAATGTGGCCGTTTACCCCTCTCAGGCCGG
CTACGCATCATCGCCTTGGTGGGCCGTTACCTCACCAACTAGCTAATG
CGCCGCAGGTCCATCCATGTTACGCCTTGATGGGCGCTTTAATATAC
TGAGCATGCGCTCTGTATACCTATCCGGTTTATAGTACCGTTTCCAGC
AGTTATCCCGGACACATGGGCTAGG
>PC.634_2 FLP3FBN01EG8AX
TTGGACCGTGTCTCAGTTCCAATGTGGGGCCCTTCCTCTCAGAACCCC
TATCCATCGAAGGCTTGGTGGGCCGTTACCCGCCAACAACCTAATGG
AACGCATCCCATCGATGACCGAAGTTCTTTAATAGTTCTACCATGCG
GAAGAACTATGCCATCGGGTATTAATCTTTCTTTTCGAAAGGCTATCCC
CGAGTCATCGGCAGGTTGGATACGTGTTACTACCCGTGCGCCGGT
>PC.354_3 FLP3FBN01EEWKD
TTGGGCCGTGTCTCAGTCCCAATGTGGCCGATCAGTCTCTTAACCTCG
CTATGCATCATTGCCTTGGTAAGCCGTACCTTACCAACTAGCTAATG
CACCGCAGGTCCATCCAAGAGTGATAGCAGAACCATCTTTCAAACCT
AGACATGCGTCTAGTGTGTTATCCGGTATTAGCATCTGTTTCCAGGT
GTTATCCAGTCTCTTGGG
...
```

BLAST against
reference tree

Reference tree
of non-redundant
full length sequences



OTU picking: example of a compute intensive process

What taxa are represented in each sample?

>PC.634_1 FLP3FBN01ELBSX

```
CTGGGCCGTGTCTCAGTCCCAATGTGGCCGTTTACCTCTCAGGCCGG
CTACGCATCATCGCCTTGGTGGGCCGTTACCTCACCAACTAGCTAATG
CGCCGCAGGTCCATCCATGTTACGCCTTGATGGGCGCTTTAATATAC
TGAGCATGCGCTCTGTATACCTATCCGGTTTGTAGCTACCGTTCCAGC
AGTTATCCCGGACACATGGGCTAGG
```

>PC.634_2 FLP3FBN01EG8AX

```
TTGGACCGTGTCTCAGTCCAATGTGGGGGCCCTTCTCTCAGAACCCC
TATCCATCGAAGGCTTGGTGGGCCGTTACCCCGCCAACAACCTAATGG
AACGCATCCCCATCGATGACCGAAGTTCTTTAATAGTTCTACCATGCG
GAAGAATAATGCCATCGGGTATTAATCTTCTTCGAAAGGCTATCCC
CGAGTCATCGGCAGGTTGGATACGTGTTACTCACCCGTGCGCCGGT
```


>PC.354_3 FLP3FBN01EEWKD

```
TTGGGCCGTGTCTCAGTCCCAATGTGGCCGATCAGTCTCTTAACTCGG
CTATGCATCATTGCCCTTGGTAAGCCGTTACCTTACCAACTAGCTAATG
CACCGCAGGTCCATCCAAGAGTGATAGCAGAACCATCTTTCAAACCTCT
AGACATGCGTCTAGTGTGTTATCCGGTATTAGCATCTGTTTCCAGGT
GTTATCCAGTCTCTTGGG
```

...

BLAST against
reference tree

Reference tree
of non-redundant
full length sequences



Clusters of “Operational Taxonomic Units” (OTUs);
Per sample hits on reference tree;
Taxonomic assignments

OTU Picking

- For 1 billion sequence reads, the initial step ran for ~116 hours on 110 processors requiring 4GB of RAM per job for workers and 64GB of RAM for master.

OTU Picking

- For 1 billion sequence reads, the initial step ran for ~116 hours on 110 processors requiring 4GB of RAM per job for workers and 64GB of RAM for the master job.
- So, on a single processor desktop with 64GB of RAM... 12760 hours or 532 days!

OTU Picking

- For 1 billion sequence reads, the initial step ran for ~116 hours on 110 processors requiring 4GB of RAM per job for workers and 64GB of RAM for the master job.
- So, on a single processor desktop with 64GB of RAM... 12760 hours or 532 days!
- One HiSeq2000 generates this data in a week!

Cloud computing

- Implemented on a cluster (or grid), but compute power is rented as a service to support arbitrary applications.

Maintaining hardware is expensive

- Temperature (redundant cooling systems)
- Redundant network connections
- Hardware maintenance (e.g., replacing hard drives)
- Fire suppression
- Back-up power
- System administrator (\$\$)

Pay-as-you-go compute power

- Public clouds (e.g., Amazon) rent compute resources
- Log in, boot virtual machine image, run analyses, and terminate instance.
- Cheaper for many tasks than buying, maintaining, and supporting a compute cluster.

Types of cloud offerings

- Applications/SaaS (e.g., Google Docs, gmail, Dropbox, iCloud)
- Computing platform/PaaS (e.g., Google App Engine)
- Raw compute resources/IaaS (e.g., Amazon Elastic Compute Cloud (EC2))

Cloud computing options

- Amazon Elastic Compute Cloud (EC2)
- Magellan – Argonne's DOE Cloud Computing
- Data Intensive Academic Grid (DIAG) –
Institute for Genome Sciences (IGS), University
of Maryland School of Medicine (UMSOM)

Interacting with the Amazon Cloud

- Boot virtual machine image via web interface (or a third-party tool like StarCluster).
- Log in and work via terminal (or via web interface with IPython Notebook)
- Move data back and forth via sftp/scp or a graphical sftp client (e.g., Cyberduck [free/cross-platform])

Virtual machines

- A “guest” operating system running within a “host” operating system
- A software implementation of a computer, that operates like a physical computer.
- A developer can create a virtual machine *image* which contains their tools pre-installed. Users can then *instantiate* that image to work with those tools.

Benefits that virtual machines offer bioinformatics

- Reproducibility: can publish protocols with a virtual machine instance id.
- Updates are burden of developer, not user.
- Coupled with cloud computing, it's the perfect model for users with sporadic compute needs.

EC2 costs: www.ec2instances.info

| Name | Memory | Compute Units | Storage | Platform | I/O Perf | Max IPs | API Name | Linux cost | Windows cost |
|---------------------------------------|----------|-----------------------------|---------------------------|-----------|-----------|---------|-------------|------------------|------------------|
| Standard Small | 1.70 GB | 1 (1 core x 1 unit) | 160 GB | 32/64-bit | Moderate | 8 | m1.small | \$0.08 per hour | \$0.115 per hour |
| Standard Medium | 3.75 GB | 2 (1 core x 2 units) | 410 GB | 32/64-bit | Moderate | 12 | m1.medium | \$0.16 per hour | \$0.23 per hour |
| Standard Large | 7.50 GB | 4 (2 cores x 2 units) | 850 GB (2 x 420 GB) | 64-bit | High | 30 | m1.large | \$0.32 per hour | \$0.46 per hour |
| Standard Extra Large | 15.00 GB | 8 (4 cores x 2 units) | 1690 GB (4 x 420 GB) | 64-bit | High | 60 | m1.xlarge | \$0.64 per hour | \$0.92 per hour |
| Micro | 0.60 GB | 2 (only for short bursts) | 0 GB (EBS only) | 32/64-bit | Low | 1 | t1.micro | \$0.02 per hour | \$0.03 per hour |
| High-Memory Extra Large | 17.10 GB | 6.5 (2 cores x 3.25 units) | 420 GB | 64-bit | Moderate | 60 | m2.xlarge | \$0.45 per hour | \$0.57 per hour |
| High-Memory Double Extra Large | 34.20 GB | 13 (4 cores x 3.25 units) | 850 GB | 64-bit | High | 120 | m2.2xlarge | \$0.90 per hour | \$1.14 per hour |
| High-Memory Quadruple Extra Large | 68.40 GB | 26 (8 cores x 3.25 units) | 1690 GB (2 x 840 GB) | 64-bit | High | 240 | m2.4xlarge | \$1.80 per hour | \$2.28 per hour |
| High-CPU Medium | 1.70 GB | 5 (2 cores x 2.5 units) | 350 GB | 32/64-bit | Moderate | 12 | c1.medium | \$0.165 per hour | \$0.285 per hour |
| High-CPU Extra Large | 7.00 GB | 20 (8 cores x 2.5 units) | 1690 GB (4 x 420 GB) | 64-bit | High | 60 | c1.xlarge | \$0.66 per hour | \$1.14 per hour |
| Cluster Compute Quadruple Extra Large | 23.00 GB | 33.5 (2 x Intel Xeon X5570) | 1690 GB (2 x 840 GB) | 64-bit | Very High | 1 | cc1.4xlarge | \$1.30 per hour | \$1.61 per hour |
| Cluster Compute Eight Extra Large | 60.50 GB | 88 (2 x Intel Xeon E5-2670) | 3370 GB (4 x 840 GB) | 64-bit | Very High | 240 | cc2.8xlarge | \$2.40 per hour | \$2.97 per hour |
| Cluster GPU Quadruple Extra Large | 22.00 GB | 33.5 (2 x Intel Xeon X5570) | 1690 GB (2 x 840 GB) | 64-bit | Very High | 1 | cg1.4xlarge | \$2.10 per hour | \$2.60 per hour |
| High I/O Quadruple Extra Large | 60.50 GB | 35 (2 x Intel Xeon X5570) | 2048 GB (2 x 1024 GB SSD) | 64-bit | Very High | 1 | hi1.4xlarge | \$3.10 per hour | \$3.58 per hour |

QIIME virtual machine


- The QIIME package distributes an EC2 virtual machine with QIIME and its (many) dependencies pre-installed.
- Dependencies include commonly used tools like BLAST, muscle, FastTree, uclust, IPython, and a lot more. A partial list is available here:
<http://qiime.org/install/install.html>
- Latest machine identifier can always be found at:
http://qiime.org/home_static/dataFiles.html

I think there is a world market for
maybe five computers.

- Thomas Watson, IBM Founder, 1943

I think there is a world market for maybe five computers.

- Thomas Watson, IBM Founder, 1943



All figures are in units of 1000.

<http://jeremyreimer.com/postman/node/329>

The democratization of DNA sequencing



Affordable
sequencing

Cloud
computing

Open-source
software

This work is licensed under the Creative Commons Attribution 3.0 United States License. To view a copy of this license, visit

<http://creativecommons.org/licenses/by/3.0/us/> or send a letter to Creative Commons, 171 Second Street, Suite 300, San Francisco, California, 94105, USA.

Feel free to use or modify these slides, but please credit me by placing the following attribution information where you feel that it makes sense: Greg Caporaso, www.caporaso.us.